



Open in app

Get started



Published in Towards Data Science

You have **2** free member-only stories left this month. [Sign up for Medium and get an extra one](#)



Bex T.

Follow

Apr 7, 2021 · 8 min read ★ · Listen



Save



Beginner's Guide to XGBoost for Classification Problems

Utilize the hottest ML library for state-of-the-art performance



[Open in app](#)[Get started](#)

Let me introduce you to the hottest Machine Learning library in the ML community — XGBoost. In recent years, it has been the main driving force behind the algorithms that win massive ML competitions. Its speed and performance are unparalleled and it consistently outperforms any other algorithms aimed at supervised learning tasks.

The library is parallelizable which means the core algorithm can run on clusters of GPUs or even across a network of computers. This makes it feasible to solve ML tasks by training on hundreds of millions of training examples with high performance.

Originally, it was written in C++ as a command-line application. After winning a huge competition in the field of physics, it started being widely adopted by the ML community. As a result, now the library has its APIs in several other languages including Python, R, and Julia.

In this post, you will learn the fundamentals of XGBoost to solve classification tasks, an overview of the massive list of XGBoost's hyperparameters and how to tune them.

Refresher on Terminology

Before we move on to code examples of XGBoost, let's refresh on some of the terms we will be using throughout the post.

Classification task: a supervised machine learning task in which one should predict if an instance is in some category by studying the instance's features. For example, by looking at the body measurements, patient history, and glucose levels of a person, you can predict whether a person belongs to the 'Has diabetes' or 'Does not have diabetes' group.

Binary classification: One type of classification where the target instance can only belong to either one of two classes. For example, predicting whether an email is a spam or not, whether a customer purchases some product or not, etc.

Multi-class classification: Another type of classification problem where the target can



[Open in app](#)[Get started](#)

If you find yourself confused by other terminology, I have written a small ML dictionary for beginners:

Codeless Machine Learning Dictionary For Dummies

Edit description

towardsdatascience.com

How to preprocess your datasets for XGBoost

Apart from basic data cleaning operations, there are some requirements for XGBoost to achieve top performance. Mainly:

- Numeric features should be scaled
- Categorical features should be encoded

To show how these steps are done, we will be using the Rain in Australia dataset from Kaggle where we will predict whether it will rain today or not based on some weather measurements. In this section, we will focus on preprocessing by utilizing Scikit-Learn Pipelines.

```
1 import pandas as pd
2
3 rain = pd.read_csv("data/weatherAUS.csv")
4
5 >>> rain.head()
```

5001.py hosted with ❤ by GitHub

[view raw](#)

[Open in app](#)[Get started](#)

```
1 >>> rain.info()
2
3 <class 'pandas.core.frame.DataFrame'>
4 RangeIndex: 145460 entries, 0 to 145459
5 Data columns (total 23 columns):
6 #   Column          Non-Null Count  Dtype
7 ---  -
8 0   Date            145460 non-null object
9 1   Location        145460 non-null object
10 2   MinTemp         143975 non-null float64
11 3   MaxTemp         144199 non-null float64
12 4   Rainfall        142199 non-null float64
13 5   Evaporation     82670 non-null float64
14 6   Sunshine        75625 non-null float64
15 7   WindGustDir     135134 non-null object
16 8   WindGustSpeed   135197 non-null float64
17 9   WindDir9am      134894 non-null object
18 10  WindDir3pm      141232 non-null object
19 11  WindSpeed9am     143693 non-null float64
20 12  WindSpeed3pm     142398 non-null float64
21 13  Humidity9am      142806 non-null float64
22 14  Humidity3pm      140953 non-null float64
23 15  Pressure9am      130395 non-null float64
24 16  Pressure3pm      130432 non-null float64
25 17  Cloud9am         89572 non-null float64
26 18  Cloud3pm         86102 non-null float64
27 19  Temp9am          143693 non-null float64
28 20  Temp3pm          141851 non-null float64
29 21  RainToday        142199 non-null object
30 22  RainTomorrow     142193 non-null object
31 dtypes: float64(16), object(7)
32 memory usage: 25.5+ MB
```

5002.py hosted with ❤ by GitHub

[view raw](#)

The dataset contains weather measures of 10 years from multiple weather stations in Australia. You can either predict whether it will rain tomorrow or today, so there are two targets in the dataset named `RainToday`, `RainTomorrow`.



[Open in app](#)[Get started](#)

```
cols_to_drop = ["Date", "Location", "RainTomorrow", "Rainfall"]

rain.drop(cols_to_drop, axis=1, inplace=True)
```

Dropping the `Rainfall` column is a must because it records the amount of rain in millimeters.

Next, let's deal with missing values starting by looking at their proportions in each column:

```
1  missing_props = rain.isna().mean(axis=0)
2
3  >>> missing_props
4
5  MinTemp      0.010209
6  MaxTemp      0.008669
7  Evaporation  0.431665
8  Sunshine     0.480098
9  WindGustDir   0.070989
10 WindGustSpeed 0.070555
11 WindDir9am    0.072639
12 WindDir3pm    0.029066
13 WindSpeed9am  0.012148
14 WindSpeed3pm  0.021050
15 Humidity9am   0.018246
16 Humidity3pm   0.030984
17 Pressure9am   0.103568
18 Pressure3pm   0.103314
19 Cloud9am      0.384216
20 Cloud3pm      0.408071
21 Temp9am       0.012148
22 Temp3pm       0.024811
23 RainToday     0.022419
24 dtype: float64
```

5003.py hosted with ❤ by GitHub

[view raw](#)

If the proportion is higher than 40% we will drop the column:



[Open in app](#)[Get started](#)

```
4
5  Evaporation    0.431665
6  Sunshine      0.480098
7  Cloud3pm      0.408071
8  dtype: float64
```

5004.py hosted with ❤ by GitHub

[view raw](#)

Three columns contain more than 40% missing values. We will drop them:

```
1  rain.drop(over_threshold.index,
2           axis=1,
3           inplace=True)
```

5005.py hosted with ❤ by GitHub

[view raw](#)

Now, before we move on to pipelines, let's divide the data into feature and target arrays beforehand:

```
1  X = rain.drop("RainToday", axis=1)
2  y = rain.RainToday
```

5006.py hosted with ❤ by GitHub

[view raw](#)

Next, there are both categorical and numeric features. We will build two separate pipelines and combine them later.

The next code examples will heavily use Sklearn-Pipelines. If you are not familiar with them, check out my separate article for the [complete guide](#) on them.

For the categorical features, we will impute the missing values with the mode of the column and encode them with One-Hot encoding:

```
1  from sklearn.impute import SimpleImputer
2  from sklearn.pipeline import Pipeline
3  from sklearn.preprocessing import OneHotEncoder
```



[Open in app](#)[Get started](#)

```
9     ]
10  )
```

5007.py hosted with ❤ by GitHub

[view raw](#)

For the numeric features, I will choose the mean as an imputer and `StandardScaler` so that the features have 0 mean and a variance of 1:

```
1  from sklearn.preprocessing import StandardScaler
2
3  numeric_pipeline = Pipeline(
4      steps=[("impute", SimpleImputer(strategy="mean")),
5             ("scale", StandardScaler())]
6  )
```

5008.py hosted with ❤ by GitHub

[view raw](#)

Finally, we will combine the two pipelines with a column transformer. To specify which columns the pipelines are designed for, we should first isolate the categorical and numeric feature names:

```
1  cat_cols = X.select_dtypes(exclude="number").columns
2  num_cols = X.select_dtypes(include="number").columns
```

5009.py hosted with ❤ by GitHub

[view raw](#)

Next, we will input these along with their corresponding pipelines into a `ColumnTransformer` instance:

```
1  from sklearn.compose import ColumnTransformer
2
3  full_processor = ColumnTransformer(
4      transformers=[
5          ("numeric", numeric_pipeline, num_cols),
6          ("categorical", categorical_pipeline, cat_cols),
7      ]
9  )
```



[Open in app](#)[Get started](#)

The full pipeline is finally ready. The only thing missing is the XGBoost classifier, which we will add in the next section.

An Example of XGBoost For a Classification Problem

To get started with `xgboost`, just install it either with `pip` or `conda`:

```
# pip
pip install xgboost

# conda
conda install -c conda-forge xgboost
```

After installation, you can import it under its standard alias — `xgb`. For classification problems, the library provides `XGBClassifier` class:

```
1  import xgboost as xgb
2
3  xgb_cl = xgb.XGBClassifier()
4
5  >>> print(type(xgb_cl))
6
7  <class 'xgboost.sklearn.XGBClassifier'>
```

5011.py hosted with ❤ by GitHub

[view raw](#)

Fortunately, the classifier follows the familiar fit-predict pattern of `sklearn` meaning we can freely use it as any `sklearn` model.

Before we train the classifier, let's preprocess the data and divide it into train and test sets:

```
1  # Apply preprocessing
```



[Open in app](#)[Get started](#)

```
7 from sklearn.model_selection import train_test_split
8
9 X_train, X_test, y_train, y_test = train_test_split(
10     X_processed, y_processed, stratify=y_processed, random_state=1121218
11 )
```

5012.py hosted with ❤ by GitHub

[view raw](#)

Since the target contains `NaN`, I imputed it by hand. Also, it is important to pass `y_processed` to `stratify` so that the split contains the same proportion of categories in both sets.

Now, we fit the classifier with default parameters and evaluate its performance:

```
1 from sklearn.metrics import accuracy_score
2
3 # Init classifier
4 xgb_cl = xgb.XGBClassifier()
5
6 # Fit
7 xgb_cl.fit(X_train, y_train)
8
9 # Predict
10 preds = xgb_cl.predict(X_test)
11
12 # Score
13 >>> accuracy_score(y_test, preds)
14 0.8507080984463082
```

5013.py hosted with ❤ by GitHub

[view raw](#)

Even with default parameters, we got an 85% accuracy which is reasonably good. In the next sections, we will try to improve the model even further by using `GridSearchCV` offered by Scikit-learn.



[Open in app](#)[Get started](#)

Unlike many other algorithms, XGBoost is an ensemble learning algorithm meaning that it combines the results of many models, called **base learners** to make a prediction.

Just like in Random Forests, XGBoost uses Decision Trees as base learners:

Visualizing a Decision Tree of Rain

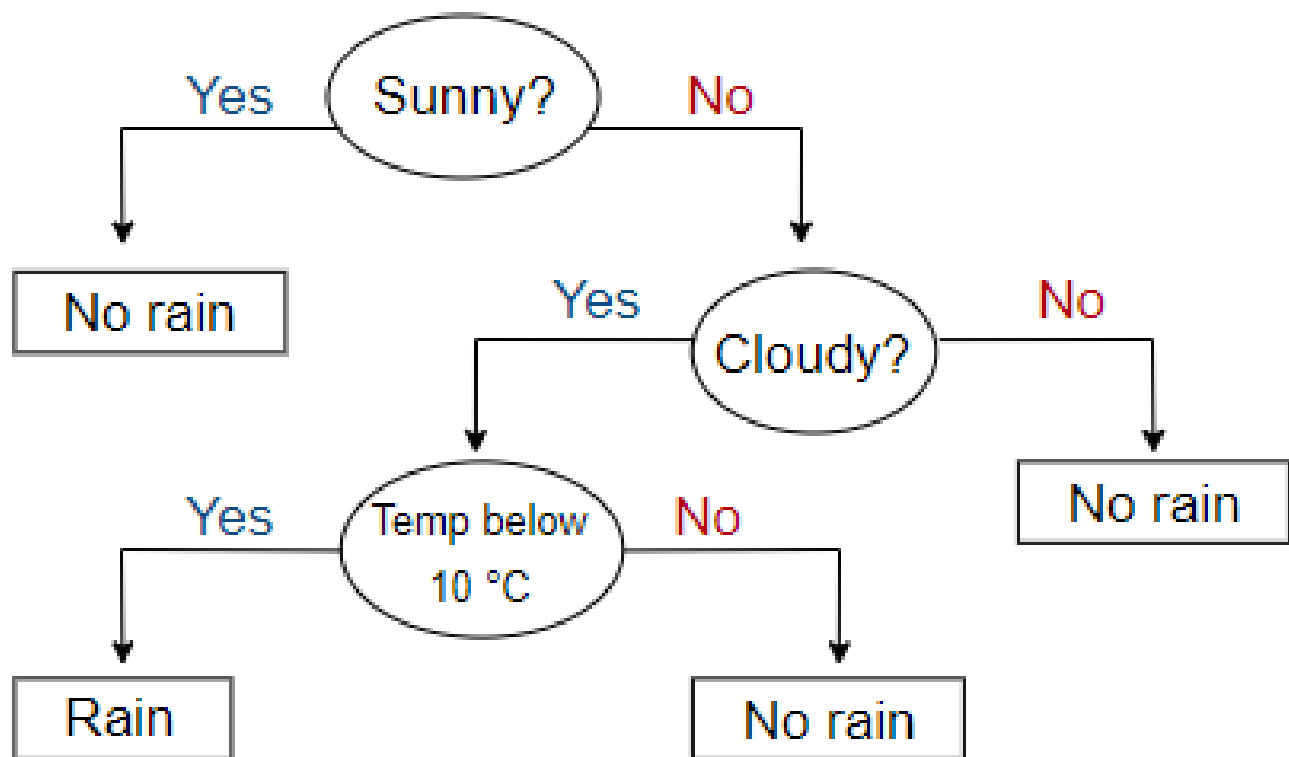


Image by the author. Decision tree to predict rain

An example of a decision tree can be seen above. In each decision node (circles), there is a single question that is being asked with only two possible answers. At the bottom of each tree, there is a single decision (rectangles). In the above tree, the first question is whether it is sunny or not. If yes, you immediately decide that it is not going to rain. If otherwise, you continue to ask more binary (yes/no) questions that ultimately will lead to some decision at the last “leaf” (rectangle)



[Open in app](#)[Get started](#)

unseen data.

However, the trees used by XGBoost are a bit different than traditional decision trees. They are called CART trees (Classification and Regression trees) and instead of containing a single decision in each “leaf” node, they contain real-value scores of whether an instance belongs to a group. After the tree reaches max depth, the decision can be made by converting the scores into categories using a certain threshold.

I am in no way an expert when it comes to the internals of XGBoost. That’s why I recommend you to check out [this awesome YouTube playlist](#) entirely on XGBoost and [another one](#) solely aimed at Gradient Boosting which I did not mention at all.

Overview of XGBoost Classifier Hyperparameters

So far, we have been using only the default hyperparameters of the XGBoost Classifier:

```
1  >>> xgb_cl
2
3  XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
4                colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
5                importance_type='gain', interaction_constraints='',
6                learning_rate=0.300000012, max_delta_step=0, max_depth=6,
7                min_child_weight=1, missing=nan, monotone_constraints='()',
8                n_estimators=100, n_jobs=4, num_parallel_tree=1, random_state=0,
9                reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
10               tree_method='exact', validate_parameters=1, verbosity=None)
```

5014.py hosted with ❤ by GitHub

[view raw](#)

Terminology refresher: hyperparameters of a model are the settings of that model which should be provided by the user. The model itself cannot learn these from the given training data.





Open in app

Get started

1. `learning_rate` : also called *eta*, it specifies how quickly the model fits the residual errors by using additional base learners.
 - typical values: 0.01–0.2
2. `gamma`, `reg_alpha`, `reg_lambda` : these 3 parameters specify the values for 3 types of regularization done by XGBoost - minimum loss reduction to create a new split, L1 reg on leaf weights, L2 reg leaf weights respectively
 - typical values for `gamma` : 0 - 0.5 but highly dependent on the data
 - typical values for `reg_alpha` and `reg_lambda` : 0 - 1 is a good starting point but again, depends on the data
3. `max_depth` - how deep the tree's decision nodes can go. Must be a positive integer
 - typical values: 1–10
4. `subsample` - fraction of the training set that can be used to train each tree. If this value is low, it may lead to underfitting or if it is too high, it may lead to overfitting
 - typical values: 0.5–0.9
5. `colsample_bytree` - fraction of the features that can be used to train each tree. A large value means almost all features can be used to build the decision tree
 - typical values: 0.5–0.9

The above are the main hyperparameters people often tune. It is perfectly OK if you don't understand them all completely (like me) but you can refer to this [post](#) which gives a thorough overview of how each of the above parameters works and how to tune them.



[Open in app](#)[Get started](#)

separate models on the given data for each combination of hyperparameters. I won't go into detail about how GridSearch works but you can check out my separate comprehensive article on the topic:

Automatic Hyperparameter Tuning with Sklearn GridSearchCV and RandomizedSearchCV

[Edit description](#)

[towardsdatascience.com](#)

We will be tuning only a few of the parameters in two rounds because of how tuning is both computationally and time-expensive. Let's create the parameter grid for the first round:

```
1 param_grid = {
2     "max_depth": [3, 4, 5, 7],
3     "learning_rate": [0.1, 0.01, 0.05],
4     "gamma": [0, 0.25, 1],
5     "reg_lambda": [0, 1, 10],
6     "scale_pos_weight": [1, 3, 5],
7     "subsample": [0.8],
8     "colsample_bytree": [0.5],
9 }
```

5015.py hosted with ❤ by GitHub

[view raw](#)

In the grid, I fixed `subsample` and `colsample_bytree` to recommended values to speed things up and prevent overfitting.

We will import `GridSearchCV` from `sklearn.model_selection`, instantiate and fit it to our preprocessed data:

```
1 from sklearn.model_selection import GridSearchCV
```



[Open in app](#)[Get started](#)

```
7  grid_cv = GridSearchCV(xgb_cl, param_grid, n_jobs=-1, cv=3, scoring="roc_auc")
8
9  # Fit
10 _ = grid_cv.fit(X_processed, y_processed)
```

5016.py hosted with ❤ by GitHub

[view raw](#)

After an excruciatingly long time, we finally got the best params and best score:

```
1  >>> grid_cv.best_score_
2  0.853810061069393
```

5017.py hosted with ❤ by GitHub

[view raw](#)

This time, I chose `roc_auc` metric which calculates the area under the ROC (receiver operating characteristic) curve. It is one of the most popular and robust evaluation metrics for unbalanced classification problems. You can learn more about it [here](#). Let's see the best params:

```
1  >>> grid_cv.best_params_
2
3  {'gamma': 1,
4   'learning_rate': 0.1,
5   'max_depth': 7,
6   'reg_lambda': 10,
7   'scale_pos_weight': 3}
```

5018.py hosted with ❤ by GitHub

[view raw](#)

As you can see, only `scale_pos_weight` is in the middle of its provided range. The other parameters are at the end of their ranges meaning that we have to keep exploring:

```
1  # Insert the new fixed values to the grid
2  param_grid["scale_pos_weight"] = [3]
3  param_grid["subsample"] = [0.8]
4  param_grid["colsample_bytree"] = [0.5]
```



[Open in app](#)[Get started](#)

```
10 param_grid["learning_rate"] = [0.3, 0.5, 0.7, 1]
```

5019.py hosted with ❤ by GitHub

[view raw](#)

We will fit a new GridSearch object to the data with the updated param grid and see if we got an improvement on the best score:

```
1 grid_cv_2 = GridSearchCV(xgb_cl, param_grid,
2                           cv=3, scoring="roc_auc", n_jobs=-1)
3
4 _ = grid_cv_2.fit(X_processed, y_processed)
5
6 >>> grid_cv_2.best_score_
7 0.8501002182976786
```

5020.py hosted with ❤ by GitHub

[view raw](#)

Looks like the second round of tuning resulted in a slight decrease in performance. We have got no choice but to stick with the first set of parameters which were:

```
1 >>> grid_cv.best_params_
2
3 {'gamma': 1,
4  'learning_rate': 0.1,
5  'max_depth': 7,
6  'reg_lambda': 10,
7  'scale_pos_weight': 3}
```

5021.py hosted with ❤ by GitHub

[view raw](#)

Let's create a final classifier with the above parameters:

```
1 final_cl = xgb.XGBClassifier(
2     **grid_cv.best_params_,
3     objective="binary:logistic",
4     colsample_bytree=0.5,
5     subsample=0.8
```



[Open in app](#)[Get started](#)

Finally, make predictions on the test set:

```
1 from sklearn.metrics import roc_auc_score
2
3 _ = final_cl.fit(X_train, y_train)
4
5 preds = final_cl.predict(X_test)
```

5023.py hosted with ❤ by GitHub

[view raw](#)

Conclusion


We have made it to the end of this introductory guide on XGBoost for classification problems. Even though we covered a lot, there are still many topics to explore in terms of XGBoost itself and on the topic of classification.

I strongly recommend you to check out the links I provided as additional sources to learn XGBoost and suggest reading more on how to tackle classification problems.

How To Touch Into the Heart of Matplotlib And Create Amazing Plots

Use Matplotlib like an absolute boss!

towardsdatascience.com



159



2





Open in app

Get started

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Get this newsletter

